

Erlang for Authoritative DNS

Anthony Eden
Founder of DNSimple

What is Authoritative DNS?

young people
have no respect
for authority
nowadays

<http://www.flickr.com/photos/31797858@N00/6337153123>

Where We Came From



<http://www.flickr.com/photos/8929612@N04/6021468948>

Sunday, September 22, 13

PowerDNS + Ruby backend + MySQL

What we tried first



<http://www.flickr.com/photos/18342073@N00/3546334679>

Sunday, September 22, 13

Ruby backend communicating with Clojure service talking to MySQL
Lua extension for PowerDNS + MySQL

What led to Erlang



<http://www.flickr.com/photos/30265340@N00/1364283972>

Sunday, September 22, 13

Known for parallel processing
Started as a way of learning Erlang

Benefits of Erlang



<http://www.flickr.com/photos/62276182@N00/3259154996>

Sunday, September 22, 13

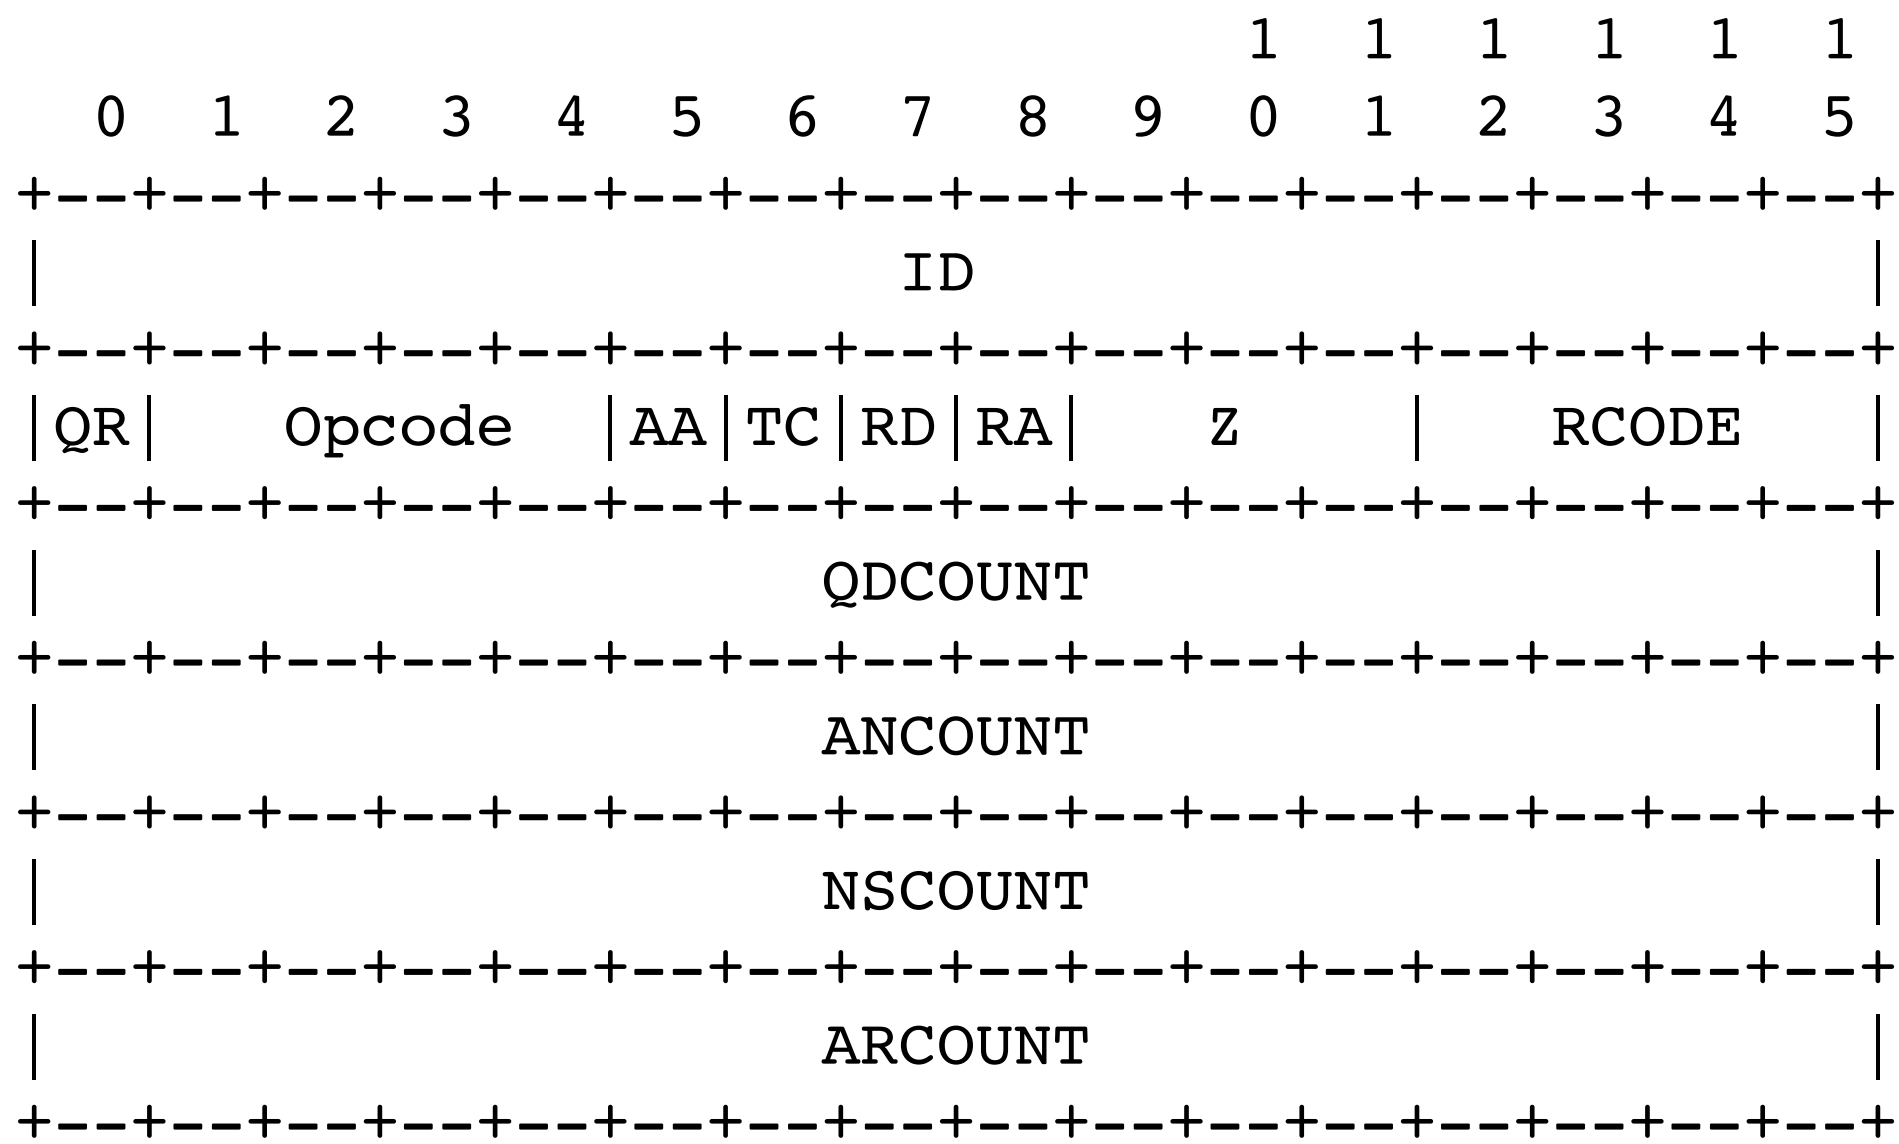
3 benefits with examples.

Binary Bit Syntax

```
decode_message(<<Id:16, QR:1, OC:4, AA:1, TC:1, RD:1, RA:  
1, 0:1, AD:1, CD:1, RC:4, QC:16, ANC:16, AUC:16, ADC:16,  
Rest/binary>> = MsgBin) ->
```



```
decode_message(<<Id:16, QR:1, OC:4, AA:1, TC:1, RD:1, RA:
1, 0:1, AD:1, CD:1, RC:4, QC:16, ANC:16, AUC:16, ADC:16,
Rest/binary>> = MsgBin) ->
```



- Sunday, September 22, 13
- Maps well to the RFC 1035 message diagram.
- ID is 16 bits (2 octets)
 - QR = Query (0) or Response (1)
 - Opcode = Four bits defining type of query (0 = standard, 1 = inverse query, 2 = server status)
 - AA = Authoritative Answer bit
 - TC = Truncation bit
 - RD = Recursion Desired bit
 - RA = Recursion Available bit
 - Z = Reserved
 - RCODE = Response code
 - QC = Question Count
 - ANC = Answer Count
 - AUC = Authority Count
 - ADC = Additional Count

Pattern Matching


```
json_record_to_erlang([Name, <<"A">>, Ttl, Data]) ->
    Ip = proplists:get_value(<<"ip">>, Data),
    case inet_parse:address(binary_to_list(Ip)) of
        {ok, Address} -> #dns_rr{
            name = Name,
            type = ?DNS_TYPE_A,
            data = #dns_rrdata_a{ip = Address},
            ttl = Ttl
        };
        {error, Reason} -> {}
    end;
```

```
json_record_to_erlang([Name, <<"CNAME">>, Ttl, Data]) ->
    #dns_rr{
        name = Name,
        type = ?DNS_TYPE_CNAME,
        data = #dns_rrdata_cname{cname =
proplists:get_value(<<"cname">>, Data)},
        ttl = Ttl
    };
```

Processes and OTP

Spawn → Init → Receive → Exit

Sunday, September 22, 13

OTP process lifecycle

Packet Cache

Sunday, September 22, 13

Maintains state.


```
-module(erldns_packet_cache).

-behavior(gen_server).

% API
-export([start_link/0, get/1, put/2, sweep/0, clear/0]).

% Gen server hooks
-export([init/1,
        handle_call/3,
        handle_cast/2,
        handle_info/2,
        terminate/2,
        code_change/3
        ]).

-define(SERVER, ?MODULE).
-define(SWEEP_INTERVAL, 1000 * 60 * 10). % Every 10 minutes

-record(state, {ttl, tref}).
```

```
% API  
-export([start_link/0, get/1, put/2, sweep/0, clear/0]).
```

Sunday, September 22, 13

Messages for get, put, sweep and clear.


```
%% Public API
start_link() ->
    gen_server:start_link({local, ?SERVER}, ?MODULE, [], []).

get(Question) ->
    gen_server:call(?SERVER, {get_packet, Question}).

put(Question, Response) ->
    gen_server:call(?SERVER, {set_packet, [Question, Response]}).

sweep() ->
    gen_server:cast(?SERVER, {sweep, []}).

clear() ->
    gen_server:cast(?SERVER, {clear}).
```

```
init([]) ->
    init([20]);

init([TTL]) ->
    ets:new(packet_cache, [set, named_table]),
    {ok, Tref} = timer:apply_interval(?SWEEP_INTERVAL, ?MODULE,
sweep, []),
    {ok, #state{ttl = TTL, tref = Tref}}.
```

Sunday, September 22, 13

Initialize the packet cache. If a TTL is not specified, use a default.
Use of pattern matching again.
State includes the TTL and a reference to the timer.


```
handle_call({get_packet, Question}, _From, State) ->
  case ets:lookup(packet_cache, Question) of
    [{Question, {Response, ExpiresAt}}] ->
      {_,T,_} = erlang:now(),
      case T > ExpiresAt of
        true ->
          lager:debug("Cache hit but expired"),
          {reply, {error, cache_expired}, State};
        false ->
          lager:debug("Time is ~p. Packet hit expires at ~p.", [T, ExpiresAt]),
          {reply, {ok, Response}, State}
      end;
    _ -> {reply, {error, cache_miss}, State}
  end;
```

```
handle_call({set_packet, [Question, Response]}, _From, State) ->
    {_, T, _} = erlang:now(),
    ets:insert(packet_cache, {Question, {Response, T + State#state.ttl}}),
    {reply, ok, State}.
```

```
handle_cast({sweep, []}, State) ->
  lager:debug("Sweeping packet cache"),
  {_, T, _} = erlang:now(),
  Keys = ets:select(packet_cache, [
    {{'$1', {'_', '$2'}}, [{ '<', '$2', T - 10}], ['$1']}
  ]),
  lager:debug("Found keys: ~p", [Keys]),
  lists:foreach(fun(K) -> ets:delete(packet_cache, K) end, Keys),
  {noreply, State};
```

```
handle_cast({clear}, State) ->
  ets:delete_all_objects(packet_cache),
  {noreply, State}.
```




Challenges

<http://www.flickr.com/photos/36698822@N00/1935755228>

Sunday, September 22, 13

Thinking in Erlang

Sunday, September 22, 13

Early challenge

Thinking in functions and processes.
Transforming data vs. updating data.

Operations

Sunday, September 22, 13

Later challenge

How do you deploy, test and operate Erlang systems?

Working around what Erlang isn't good at

Sunday, September 22, 13

Based on our needs:

Erlang is not suited for data manipulation on large amounts of data.

Erlang is not suited for producing large amounts of structured data.



Performance Tuning

Sunday, September 22, 13

Your mileage may vary.

Identify Bottlenecks

Sunday, September 22, 13

Measure from the outside.

Measure from the inside.

Optimize sequential paths

Sunday, September 22, 13

Amdahl's Law: The speedup of a program using multiple processors in parallel computing is limited by the time needed for the sequential fraction of the program.

Minimize the amount of time spent in single-threaded processes.
As soon as possible hand off to an independent, parallel process.
Be aware that the time to hand off to the process can be significant.

Reduce garbage collection

Sunday, September 22, 13

Copying data can cause excessive garbage collection.
Reuse processes wherever possible.

Where we ended up



<http://www.flickr.com/photos/24387390@N08/3585729893>

Sunday, September 22, 13

Packet read loop

Sunday, September 22, 13

Single process

Measured within using timer module.

Measured without using a timed test suite.

Optimized the tight loop.

Initially processed all requests in the loop, 30 to 250ms per request
then moved to poolboy, 4 ms per request
then moved to queue module, <1ms per request


```
handle_request(Socket, Host, Port, Bin, State) ->
  case queue:out(State#state.workers) of
    {{value, Worker}, Queue} ->
      gen_server:cast(Worker, {udp_query, Socket, Host, Port, Bin}),
      {noreply, State#state{workers = queue:in(Worker, Queue)}};
    {empty, _Queue} ->
      lager:info("Queue is empty, dropping packet"),
      {noreply, State}
  end.
```

```
handle_info({udp, Socket, Host, Port, Bin}, State) ->
    Response = folsom_metrics:histogram_timed_update(
        udp_handoff_histogram, ?MODULE, handle_request,
        [Socket, Host, Port, Bin, State]
    ),
    inet:setopts(State#state.socket, [{active, once}]),
    Response;
```

UDP Handoff 99th Percentile



Sunday, September 22, 13

99th percentile ranges from 35 microseconds to 60 microseconds.

17,000 – 25000 questions per second

680,000 – 1,000,000 qps across 40 nodes

Process pool

Sunday, September 22, 13

of workers is configured externally.

When the UDP handler is initialized all workers are created.

queue:out to retrieve a worker + cast + queue:in

Assumption is the worker will finish processing before the queue is exhausted, drop packet otherwise

In-memory cache of zone data

Sunday, September 22, 13

Originally we pulled from a database.

MySQL, then Postgres

Ultimately the best performance was to load the zones in memory

This presents its own challenges

Naming and caching – hard problems.

Remote zone loading from a special purpose server

Sunday, September 22, 13

Sacrifice start up time.

Minimize the transformation of zone data.

Let Golang handle production of zone data.

Current Performance Metrics

Sunday, September 22, 13

Folsom all the things.

Where we're headed

<http://www.flickr.com/photos/33118864@N00/1391641134>

Sunday, September 22, 13

SO_REUSEPORT

Sunday, September 22, 13

Multiple processes read from the same UDP port.

Metrics

Sunday, September 22, 13

Identify other high-latency paths and optimize those.
Example: start up bottleneck, which can be parallel.

Command & Control

Sunday, September 22, 13

Embrace the Erlang tools, but don't give up on other tooling.
Erlang tools include the ability to connect to and operate on running nodes.
Another option is to embed Cowboy to provide an HTTP API for command and control.



Erlang's sweet spot

Sunday, September 22, 13

Questions?

<http://www.flickr.com/photos/60368684@N00/8578067721>

Sunday, September 22, 13

Erlang for Authoritative DNS

Anthony Eden
Founder of DNSimple

<http://github.com/aetrion/erl-dns>